


Process Lifecycle and Unix Process Creation

Process Lifecycle


- To execute an instance of a program, *new process* must be created
 - Upon creation, process becomes active or *ready* for execution



- Multiple processes can be simultaneously *ready*
 - E.g., browser, music player, messaging app, etc. can be simultaneously open

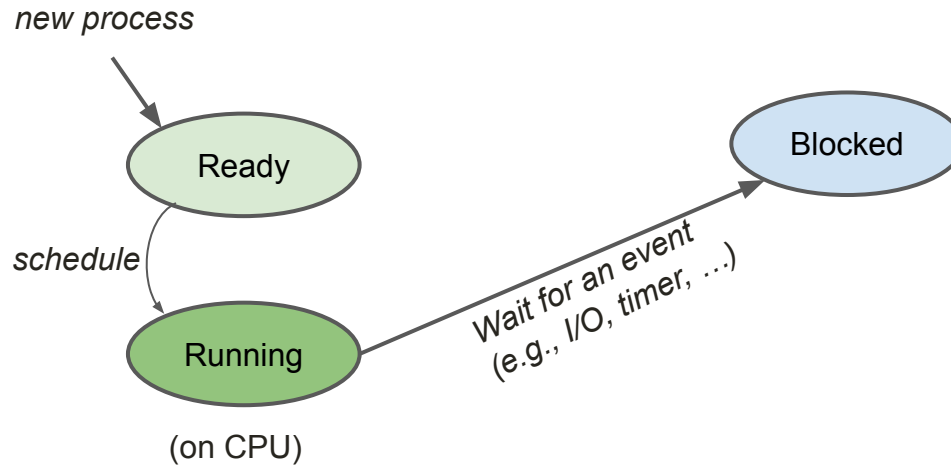
Process Lifecycle

- When CPU is free, OS chooses / schedules a ready process to *run* on the CPU
 - The chosen / scheduled process can actively execute instructions on the CPU



Process Lifecycle

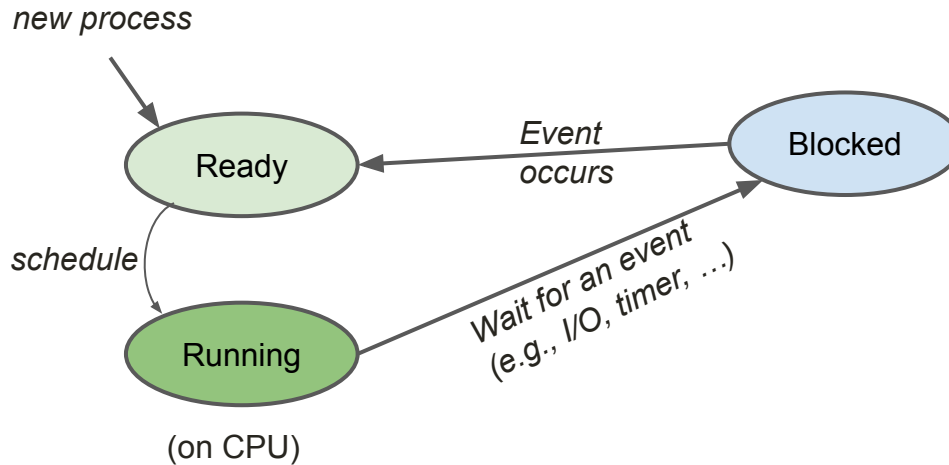
- Running process may sometimes need to pause & **wait** for some event
 - Process is put into **blocked** state so that it doesn't unnecessarily occupy the CPU



- In the meantime, OS schedules another **ready** process to **run**
 - I.e., OS switches between processes → concept of **multiprogramming**

Process Lifecycle


- Once event that **blocked** process is waiting on occurs...
 - ...process becomes **ready** again



- This **ready – running – blocked** cycle for a process can repeat

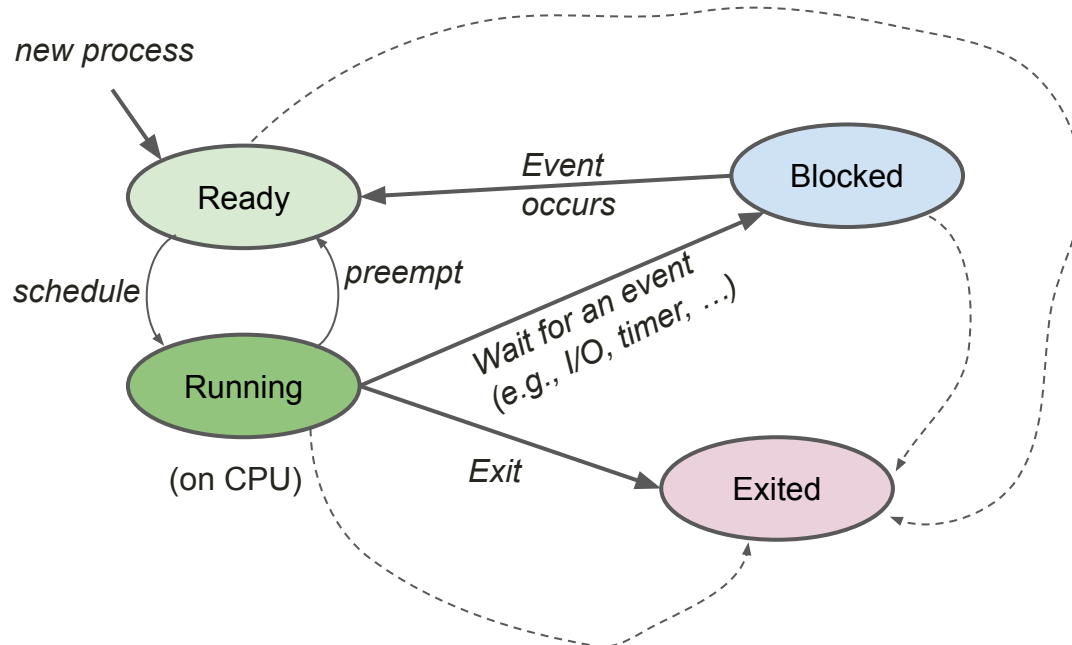
Process Lifecycle

- Once process is done, it **exits**
- Process can actually exit from any state
 - E.g., user manually kills process, process gets killed due to error, etc.



Process Lifecycle

- While process is *running*, OS may pause/*preempt* it and schedule another process (e.g., due to priority or simply for timesharing)
 - Preempted process goes back to *ready* state



- Newly scheduled processes moves to *running* state

Process Lifecycle

- Every time execution of running process is paused (due to blocking / preemption)
 - Its execution context needs to be **saved** so that it can resume later
- Every time a ready process is scheduled
 - Its execution context needs to be **loaded** / **restored** for it to run
- → **Context Switching**

To enable multiprogramming & context switching...

- ...OS must maintain information & execution context for each process
 - **Process control struct / block** used for this
 - Process ID
 - Process State (ready, running etc.)
 - Program Counter – address of next instruction to be executed
 - Registers – general purpose registers, stack pointer etc.
 - Scheduling information
 - Memory management information
 - Accounting information – time limits, etc.
 - ...


Now that we understand the process lifecycle in general, let's discuss process creation and termination in Unix-based systems

Unix process creation (*fork*)

- Existing process makes *fork* system call to create *new process*
 - First process (*init*) gets created at system start
- Creating process is the *parent* & new process is the *child*
 - Data from parent process *copied* to memory of child process
 - Memory image, environment settings, I/O handles, etc.
 - *Child* gets its own ID, process control block, scheduling info, etc.
- *Hierarchy* of parent-child relationships exist between active processes

Unix process creation (`fork()`)

- E.g.: If process A creates two new processes (B and C)...
 - ...and process C in turn creates another new process (D)



Differentiating between parent & child

- If process creation fails, **fork** returns -1 to parent & no child is created
- Successful **fork** call returns different values to parent and child
 - Fork returns child's process ID to parent process...
 - ...and returns 0 to child process
- Parent/child independently continue from execution point immediately after **fork**
 - Order of execution of the two processes may vary

```
pid_t id;

id = fork(); /* create new process */
if(id == -1) { /* error creating process */
    printf ("Error creating process\n");
} else if (id == 0) { /* only child process executes this */
    printf ("I'm a child process!\n");
} else { /* only parent process executes this */
    printf ("I just became a parent!\n");
}
```



Switching programs

- Child can *change* the program it executes
 - Invoke **exec** family of system calls to change program that it executes
 - Child *stops* executing code in parent program & *starts* executing new program

```
pid_t id;
char *argv[2];

args[0] = "echo"; /* name of new program for child to execute */
args[1] = "hello"; /* parameter to pass to the above program */
args[2] = NULL; /* end of parameter list */

id = fork(); /* create new process */
if(id == -1) { /* error creating process */
    printf ("Error creating process\n");
} else if (id == 0) { /* only child process executes this */
    execvp(args[0], args); /* child process switches to program indicated by
                           args[0] with parameters in rest of args array */
} else { /* only parent process executes this */
    printf ("I just became a parent!\n");
}
```

Process termination

- As discussed earlier, process may terminate for many reasons
 - Voluntary exit upon task completion
 - Voluntary exit due to fatal error
 - Involuntary exit due to error/bug
 - Involuntary exit due to *kill* command by OS or other process who is authorized to do this

On Unix based systems...

- Parent must be allowed to read child's *exit status* or *return value*
 - Parent process can reap children by waiting for them to terminate
 - OS provides system call for this → *wait*
- If *parent* terminates before *child*...
 - Child is now an *orphan* process
 - Orphan processes are *adopted* by *init* process
- If a *child* process terminates before *parent*
 - System will still need to keep child's *process control struct*
 - Child process becomes a *zombie* process
 - “Dead”, but not “reaped”



Waiting for child to terminate

```
pid_t id, ret;
int status;

id = fork(); /* create new process */
if(id == -1) { /* error creating process */
    printf ("Error creating process\n");
} else if (id == 0) { /* only child process executes this */
    char * args[] = {"echo", "hello", NULL};
    execvp(args[0], args); /* child process switches to program indicated by
                           args[0] with parameters in rest of args array */
} else { /* only parent process executes this */
    printf ("I just became a parent!\n");
    ret = wait(&status); /* parent process waits for child to terminate */
}
```

